

Transport Layer Protocols

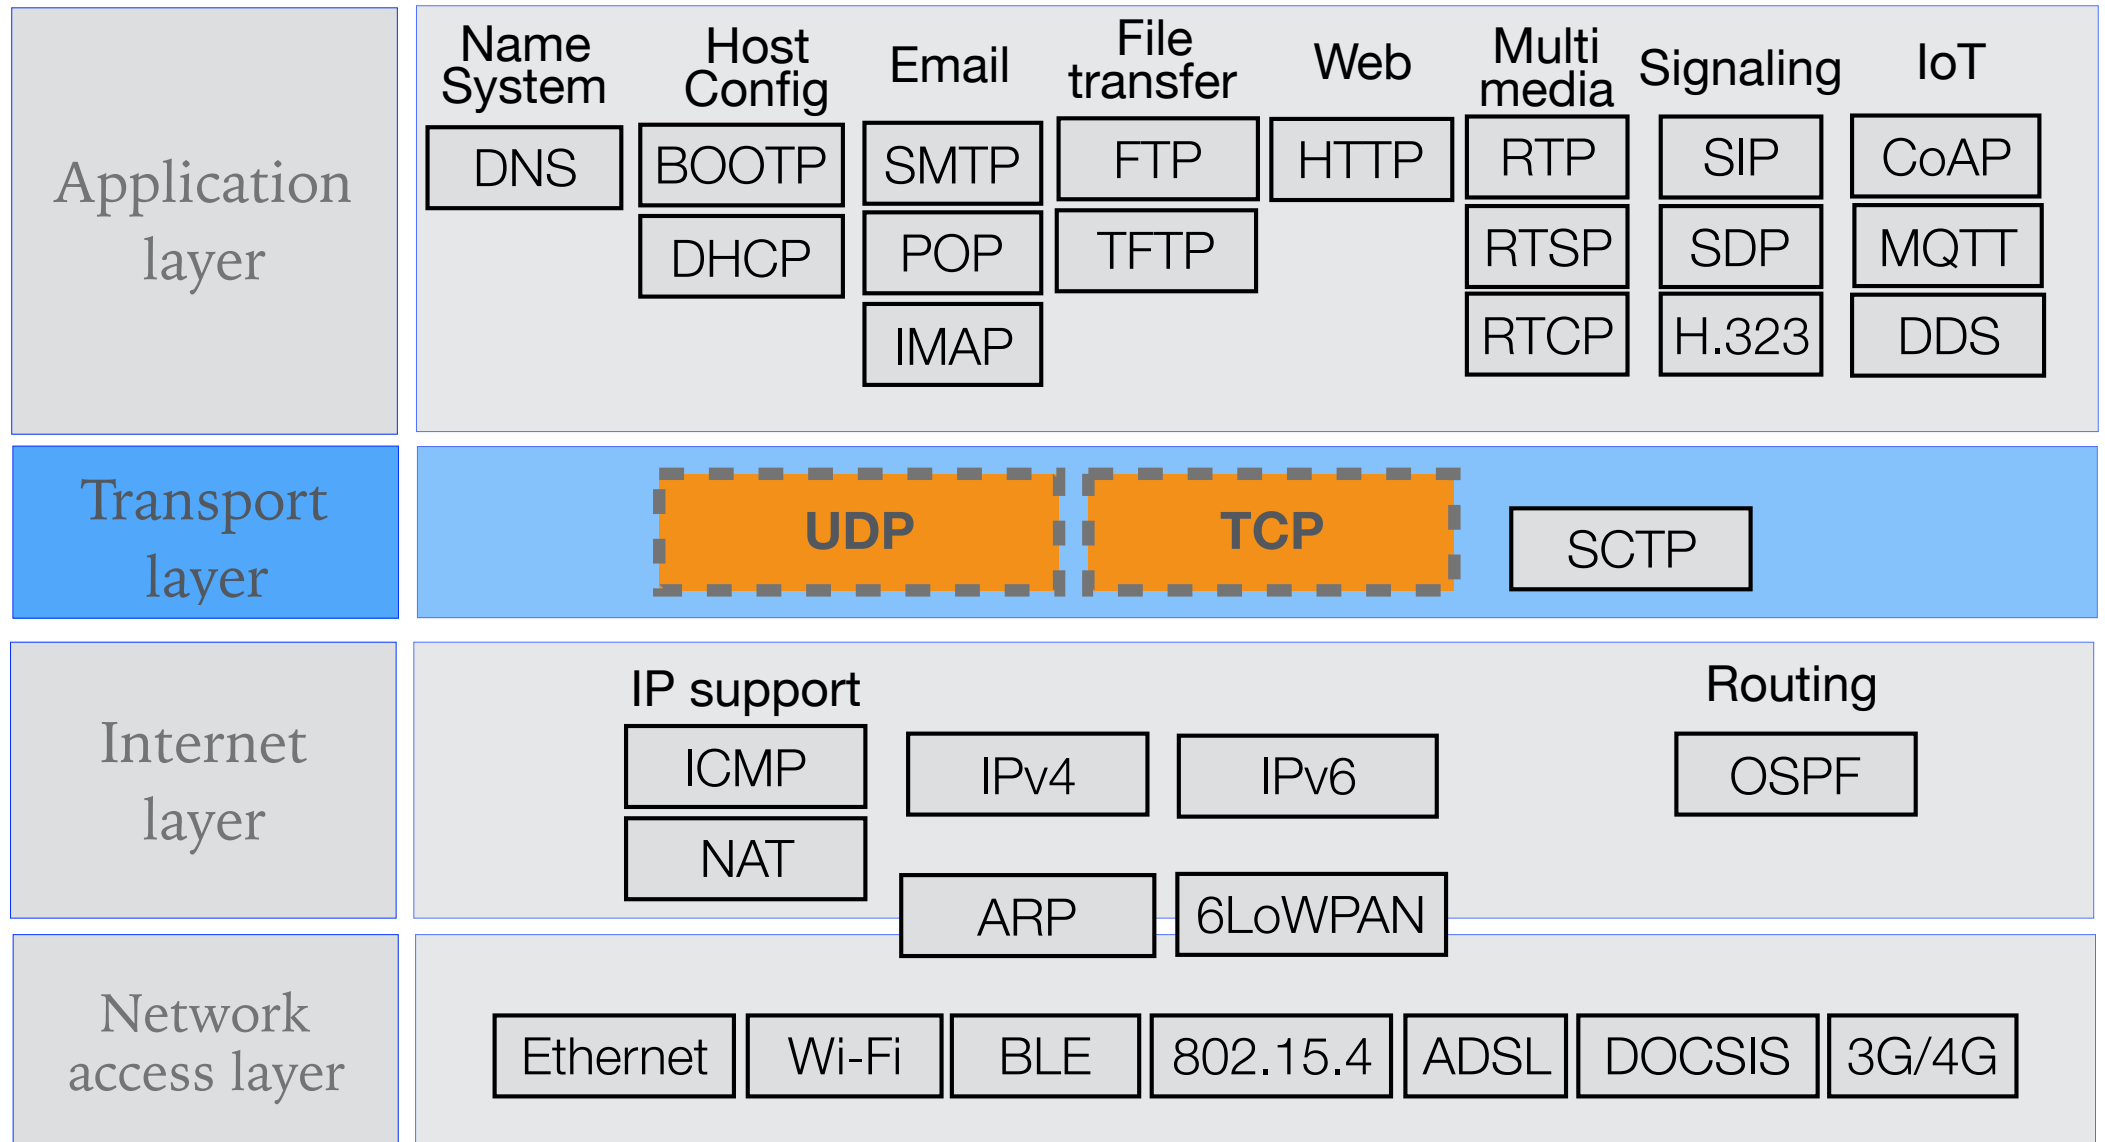
Peerapon S.

CPE 314: Computer Networks (2/61)

Topics

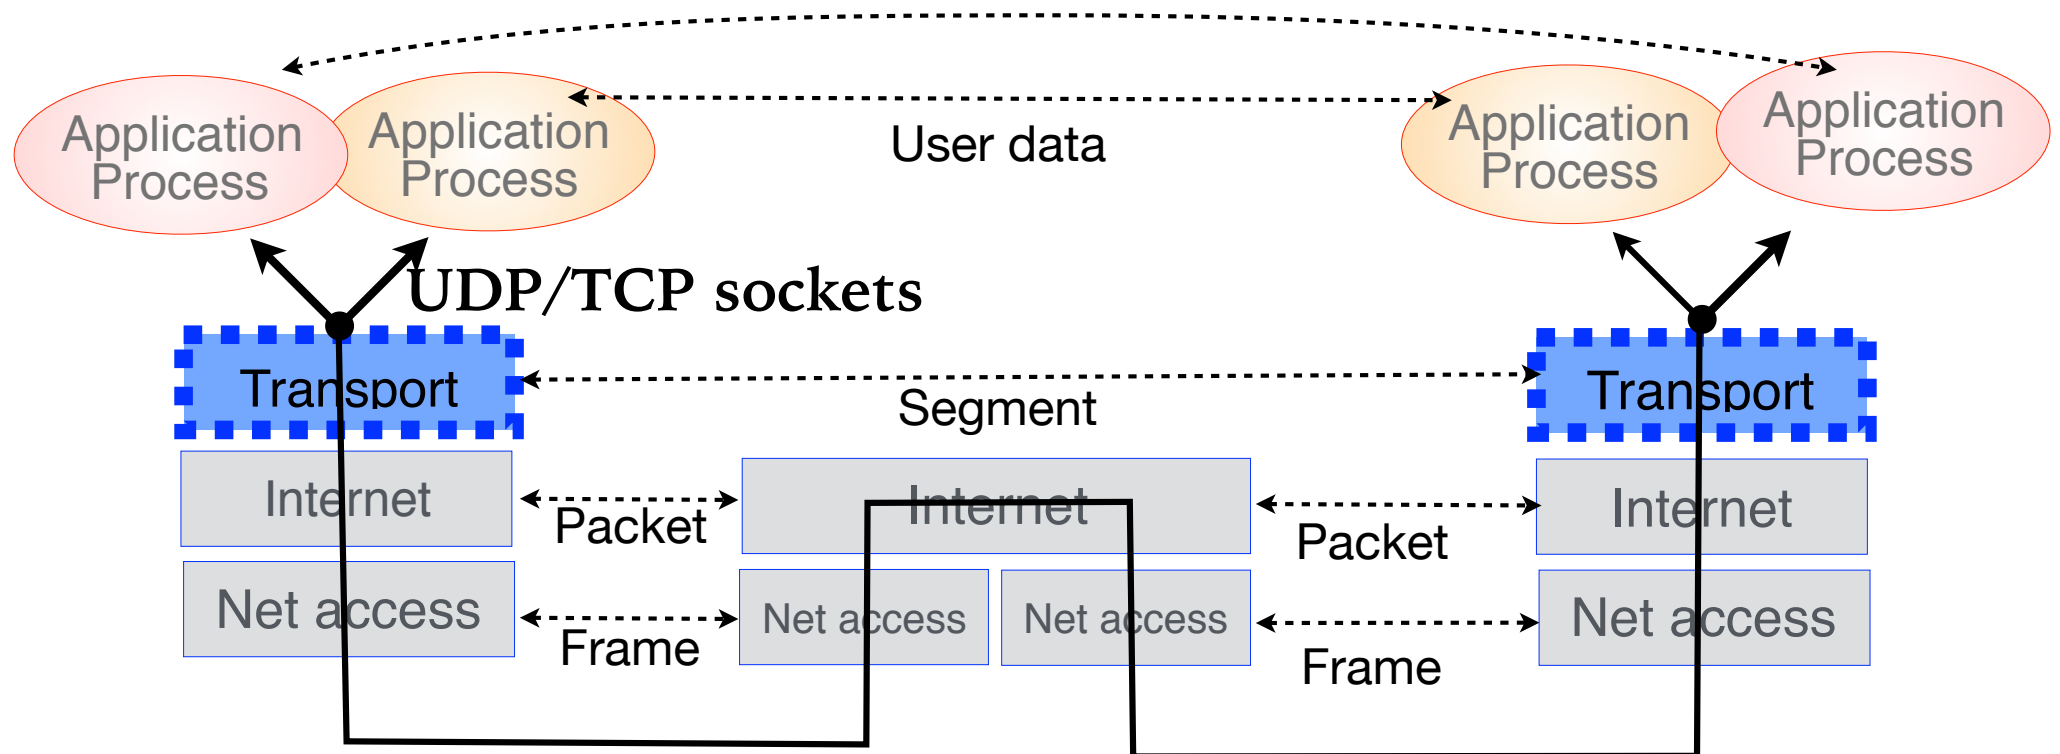
- ❑ User Datagram Protocol (UDP)
- ❑ Transport Control Protocol (TCP)
 - TCP connection management
 - TCP flow control

Protocols in TCP/IP Suite



Transport Layer (TL) Services

- Process-to-process message delivery service



Two Main Internet Transport Layer Protocols

■ UDP: User Datagram Protocol

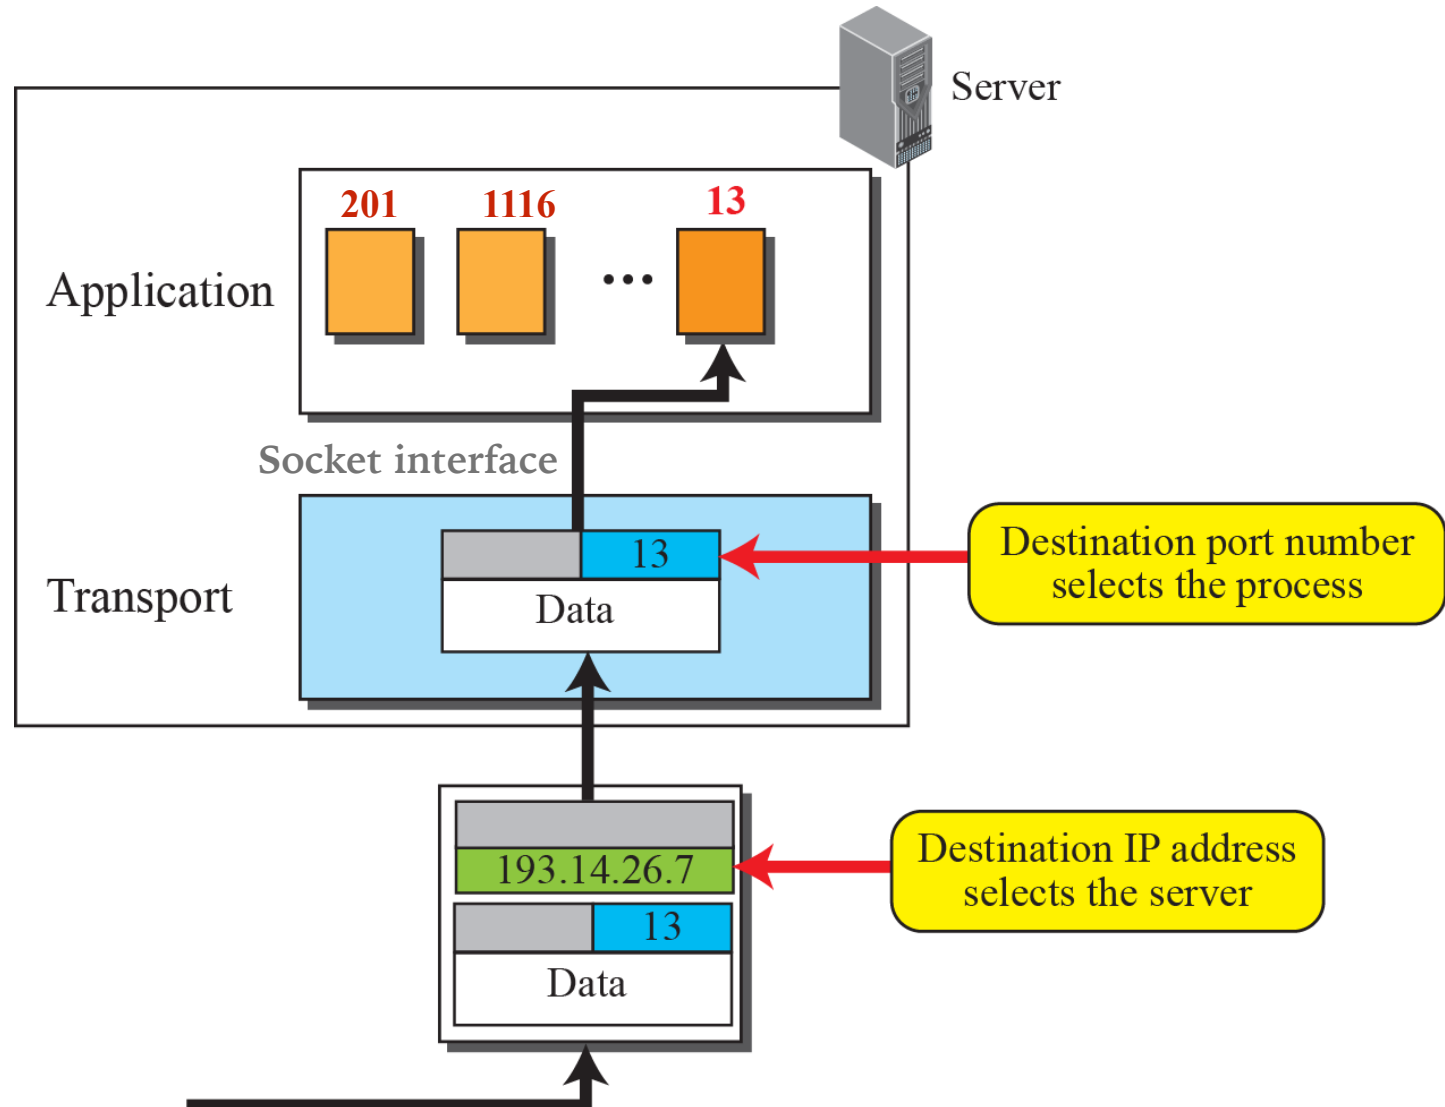
- Connection-less, Unreliable service
- Independent messages delivery

■ TCP: Transport Control Protocol

- Connection-oriented, Reliable service
- Byte-stream delivery

	Loss-free	Order	No duplicate	Throughput	Bound delay
UDP					
TCP	●	●	●		

Basic Function: Process Multiplexing



Port Range for Server and Client

- Only some port numbers can be used by your custom app.



0 - 1023	1024 - 49,151	49,152 - 65,535
----------	---------------	-----------------

Well-known
service ports

Registered ports

Ephemeral ports
(dynamic, private, temporary)

	Local port	Remote port
Client side	Random from Ephemeral range	Service or Registered port
Server side	Service or Registered port	Derived from client segment



b. Header format

- ▶ Frame 482: 215 bytes on wire (1720 bits), 215 bytes captured (1720 bits) on int
- ▶ Ethernet II, Src: Azurewav_6f:1d:03 (54:27:1e:6f:1d:03), Dst: IPv4mcast_7f:ff:f
- ▶ Internet Protocol Version 4, Src: 10.35.245.35, Dst: 239.255.255.250

▼ User Datagram Protocol, Src Port: 51082, Dst Port: 1900

Source Port: 51082

Destination Port: 1900

Length: 181

Checksum: 0xfefd [unverified]

[Checksum Status: Unverified]

[Stream index: 73]

▶ [Timestamps]

▶ Simple Service Discovery Protocol

0020	ff fa c7 8a 07 6c 00 b5 fe fd 4d 2d 53 45 41 52	...l...M-SEAR
0030	43 48 20 2a 20 48 54 54 50 2f 31 2e 31 0d 0a 48	CH * HTTP/1.1..H
0040	4f 53 54 3a 20 32 33 39 2e 32 35 35 2e 32 35 35	OST: 239 .255.255
0050	2e 32 35 30 3a 31 39 30 30 0d 0a 4d 41 4e 3a 20	.250:190 0..MAN:
0060	22 73 73 64 70 3a 64 69 73 63 6f 76 65 72 22 0d	"ssdp:discover"

Help

Close

UDP (Internet) Checksum

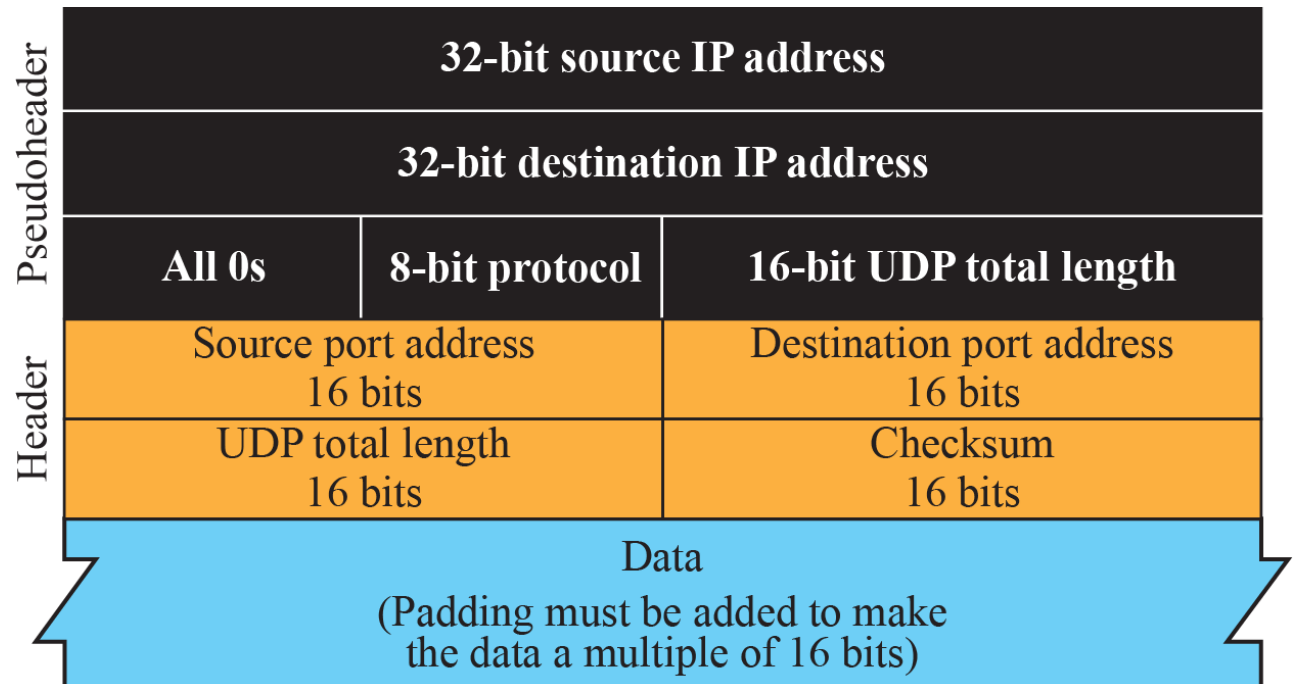
1 1 1 0 0 1 1 0 0 1 1 0 0 1 1 0

1 1 0 1 0 1 0 1 0 1 0 1 0 1 0 1

1 1 0 1 1 1 0 1 1 1 0 1 1 1 0 1

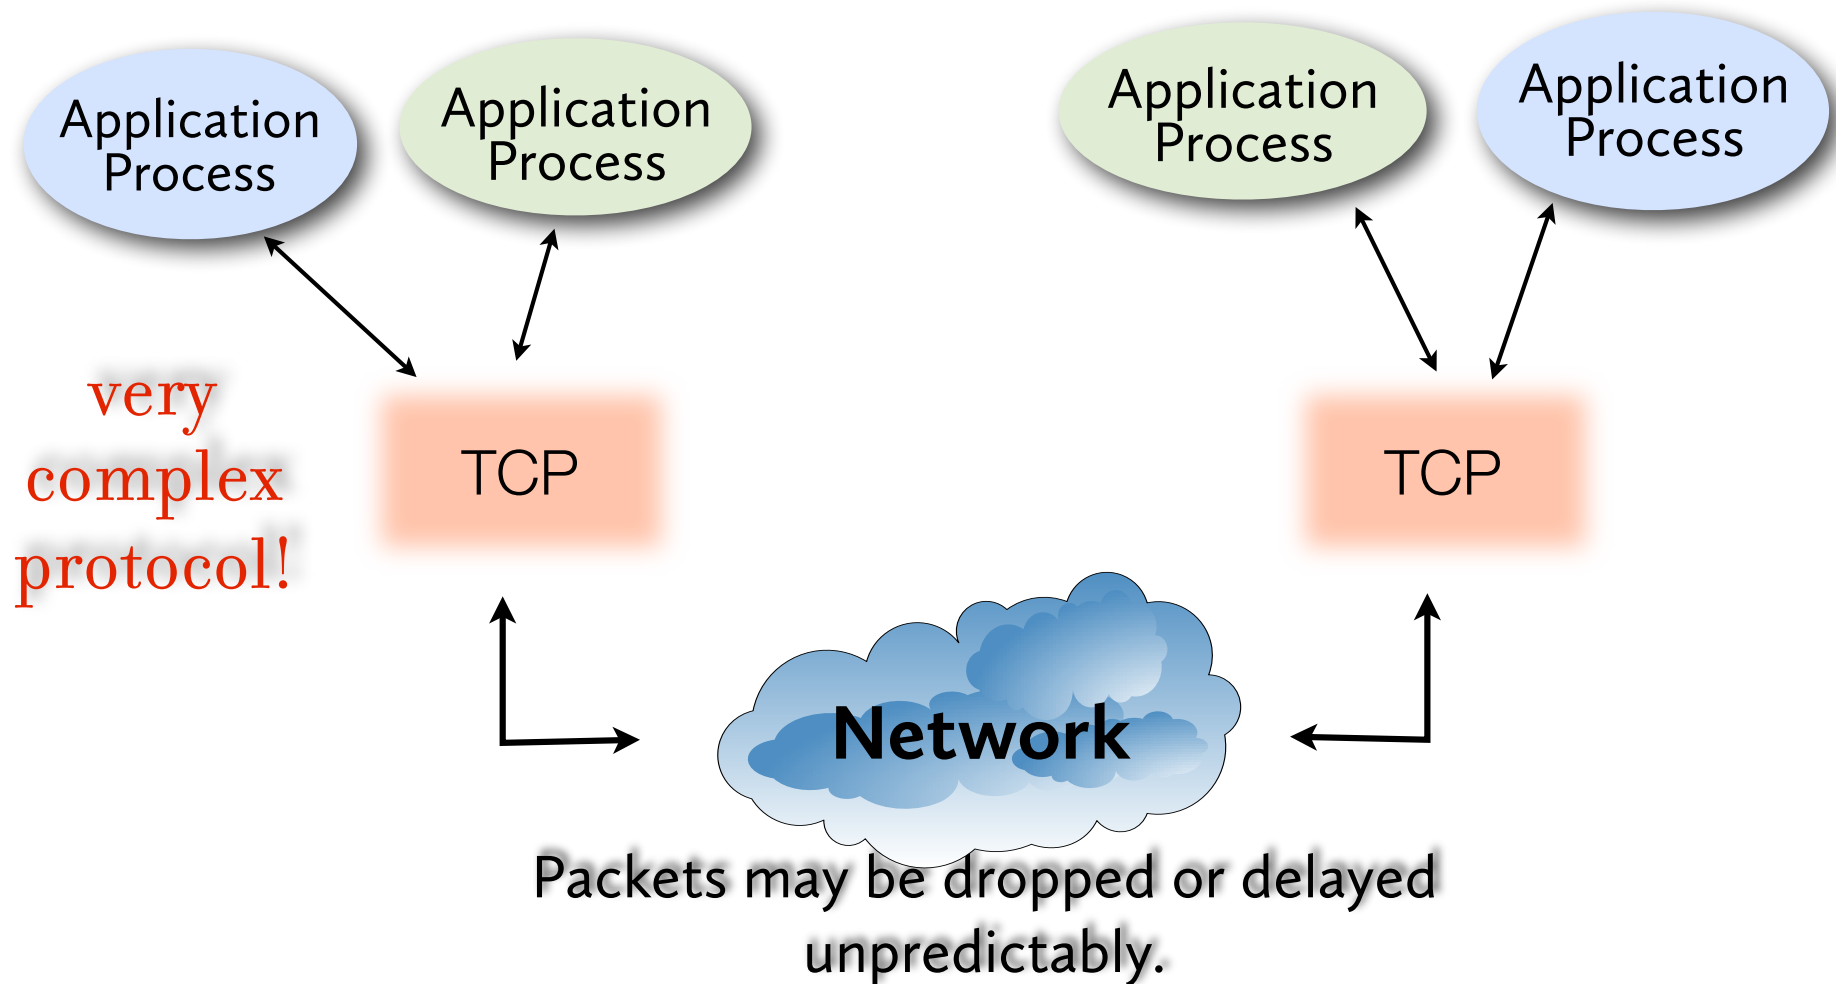
1 0 1 1 1 0 1 1 1 0 1 1 1 0 1 1

0 1 0 0 0 1 0 0 0 1 0 0 0 1 0 0

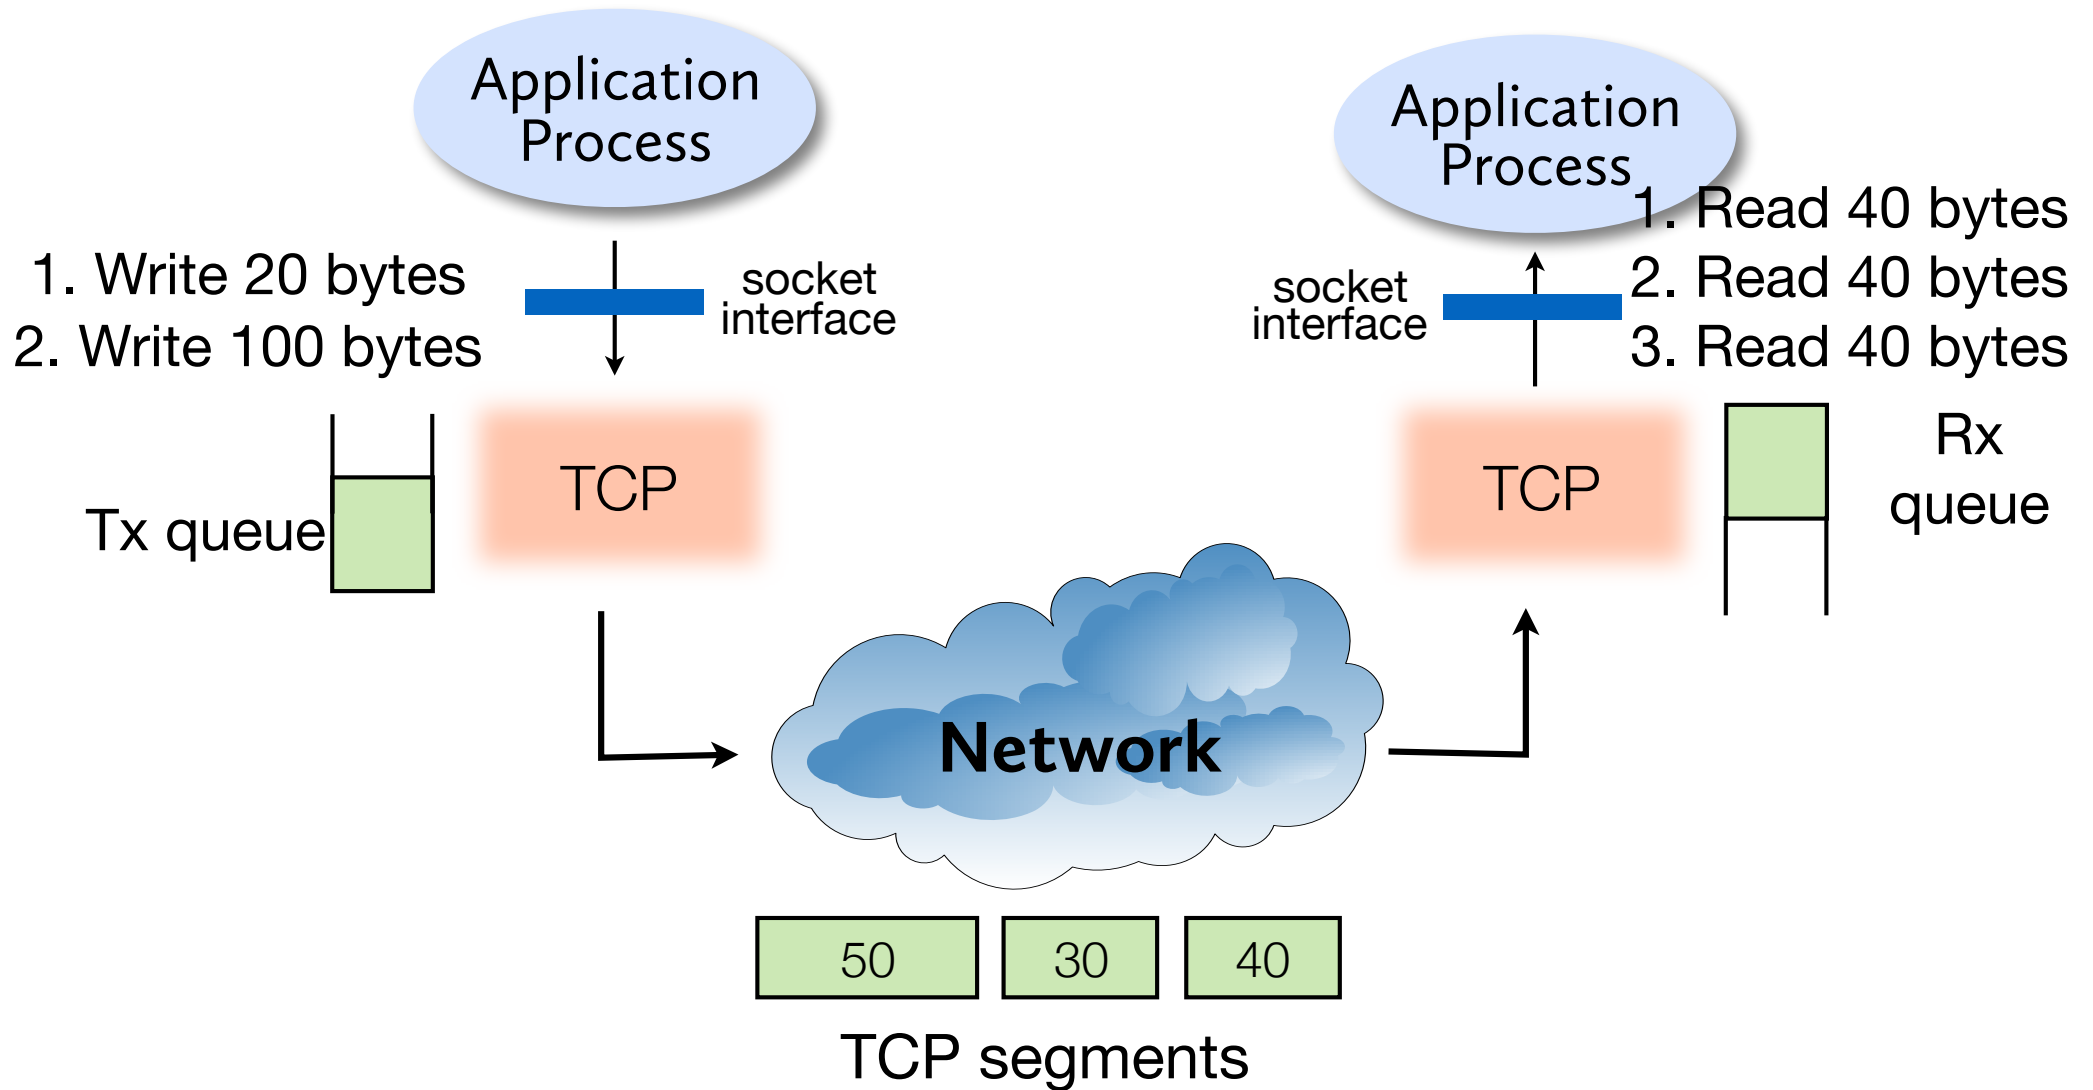


TCP: Connection-Oriented Transport


- **Reliable, full-duplex, byte-stream delivery** over an unreliable network.



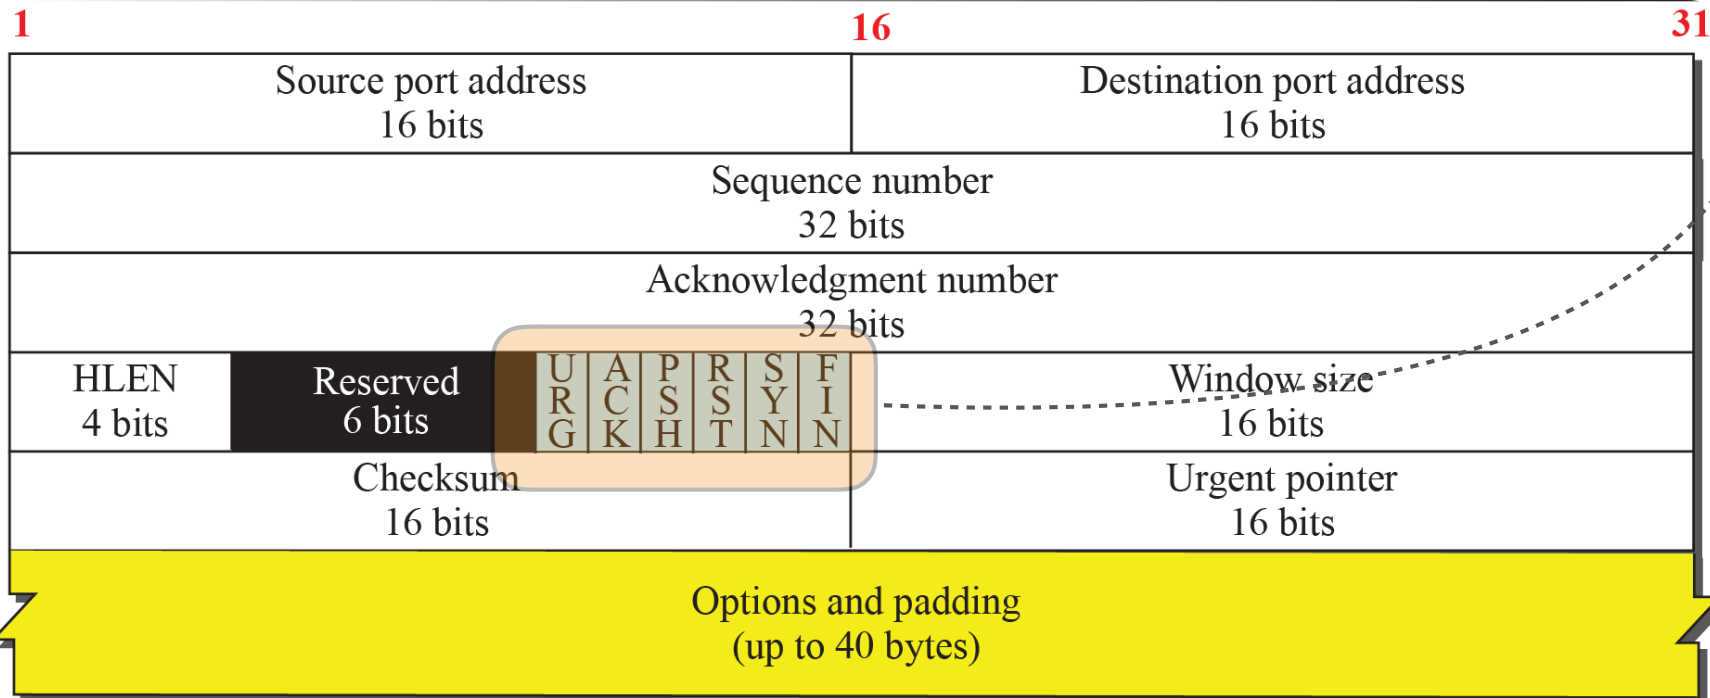
Byte-Stream Service



TCP Segment Format



URG: Urgent pointer is valid
 ACK: Acknowledgment is valid
 PSH: Request for push
 RST: Reset the connection
 SYN: Synchronize sequence numbers
 FIN: Terminate the connection



b. Header

TCP Byte Numbering and Sequence No.

- ❑ Number bytes sequentially. Why?
- ❑ Segment seq. no. = Seq. no of 1st byte in segment
- ❑ Ex: Send 5000-byte file in 5 segments, starting with byte number 10001.

Segment 1	→	Sequence Number:	10,001	Range:	10,001	to	11,000
Segment 2	→	Sequence Number:	11,001	Range:	11,001	to	12,000
Segment 3	→	Sequence Number:	12,001	Range:	12,001	to	13,000
Segment 4	→	Sequence Number:	13,001	Range:	13,001	to	14,000
Segment 5	→	Sequence Number:	14,001	Range:	14,001	to	15,000

- ▶ Frame 469: 584 bytes on wire (4672 bits), 584 bytes captured (4672 bits) on interface
- ▶ Ethernet II, Src: HewlettP_2c:26:75 (b8:af:67:2c:26:75), Dst: Apple_33:c9:54 (b8:e8:56)
- ▶ Internet Protocol Version 4, Src: 17.248.154.110, Dst: 10.35.244.167

▼ Transmission Control Protocol, Src Port: 443, Dst Port: 62281, Seq: 998675364, Ack: 13

Source Port: 443

Destination Port: 62281

[Stream index: 19]

[TCP Segment Len: 518]

Sequence number: 998675364

[Next sequence number: 998675882]

Acknowledgment number: 1357145119

1000 = Header Length: 32 bytes (8)

- ▶ Flags: 0x018 (PSH, ACK)

Window size value: 1075

[Calculated window size: 34400]

[Window size scaling factor: 32]

Checksum: 0x544d [unverified]

[Checksum Status: Unverified]

Urgent pointer: 0

- ▶ Options: (12 bytes), No-Operation (NOP), No-Operation (NOP), Timestamps

- ▶ [SEQ/ACK analysis]

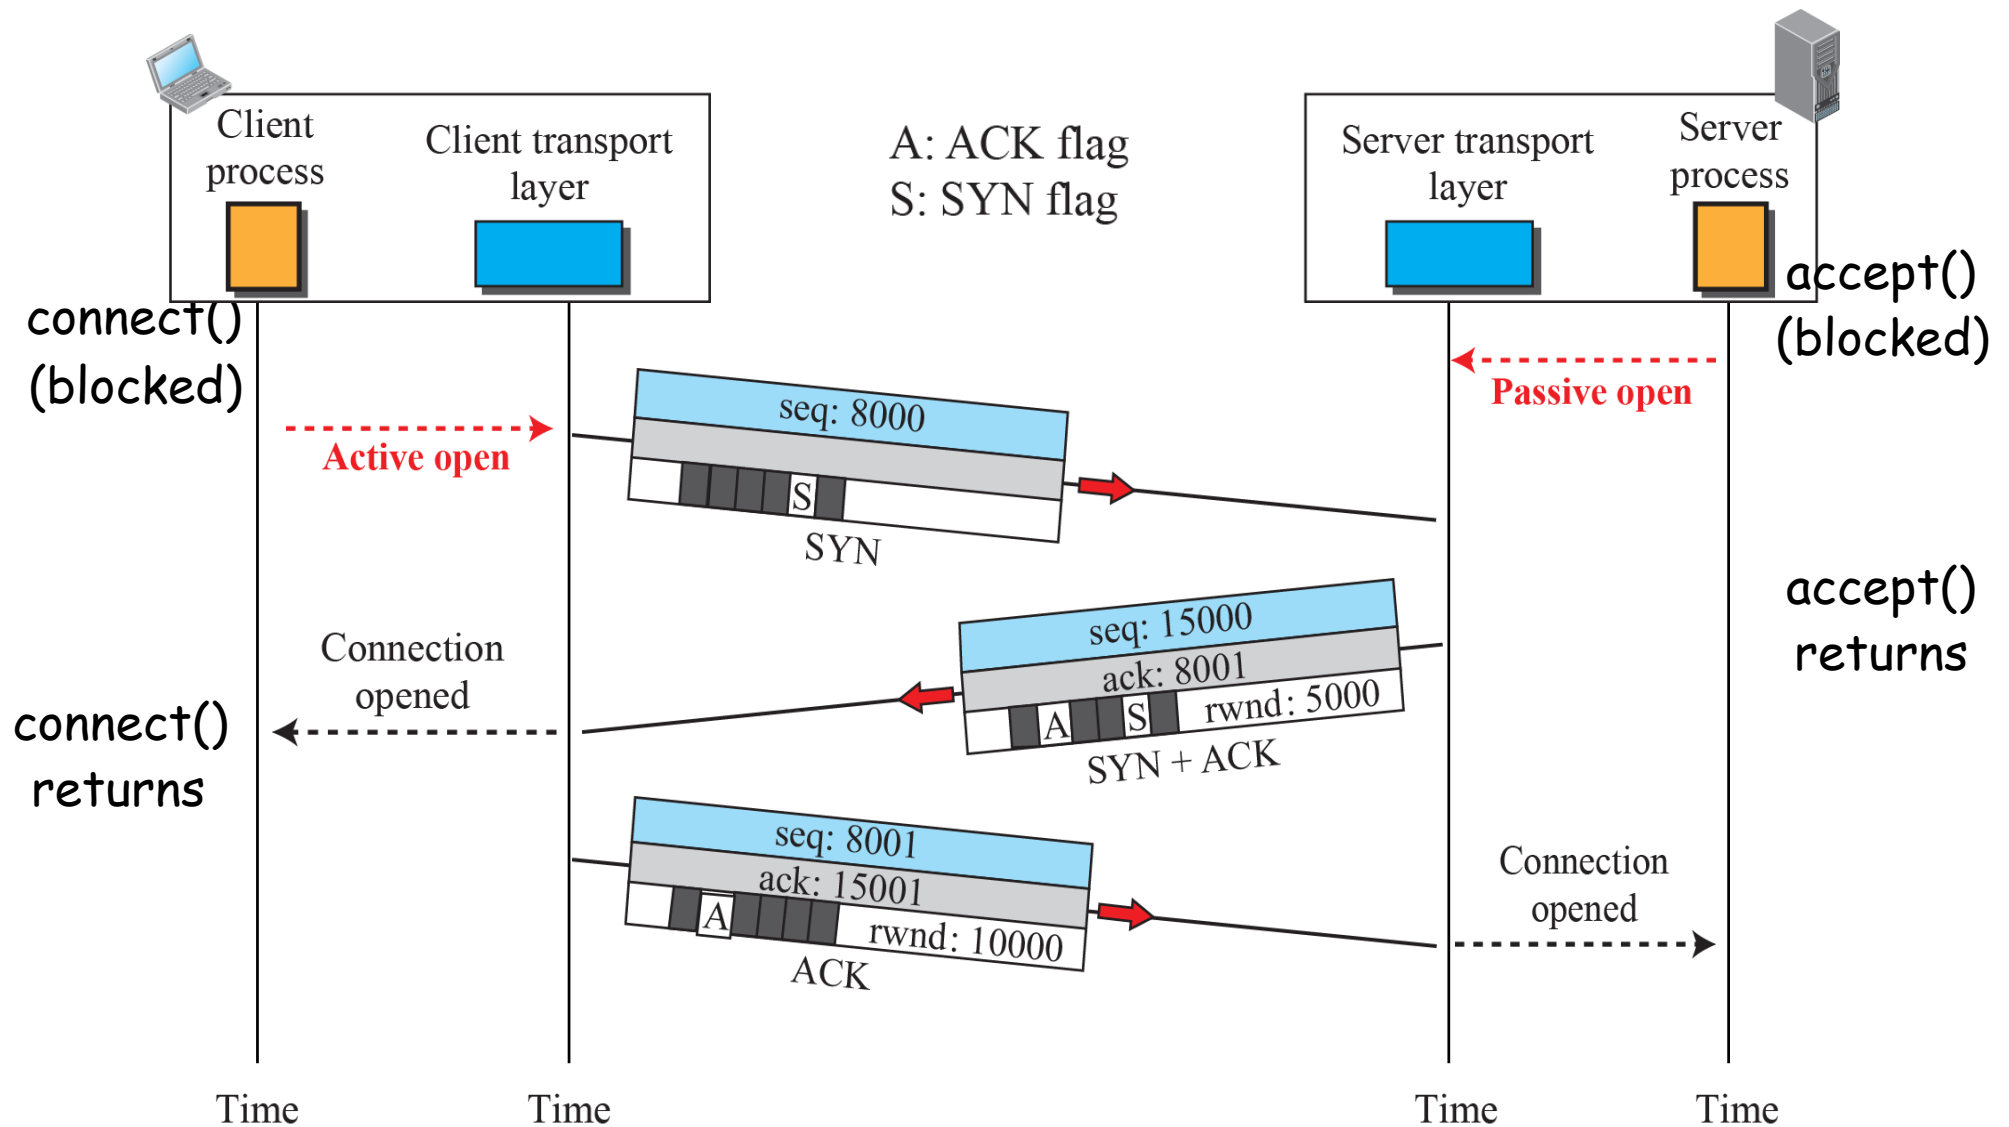
- ▶ [Timestamps]

TCP payload (518 bytes)

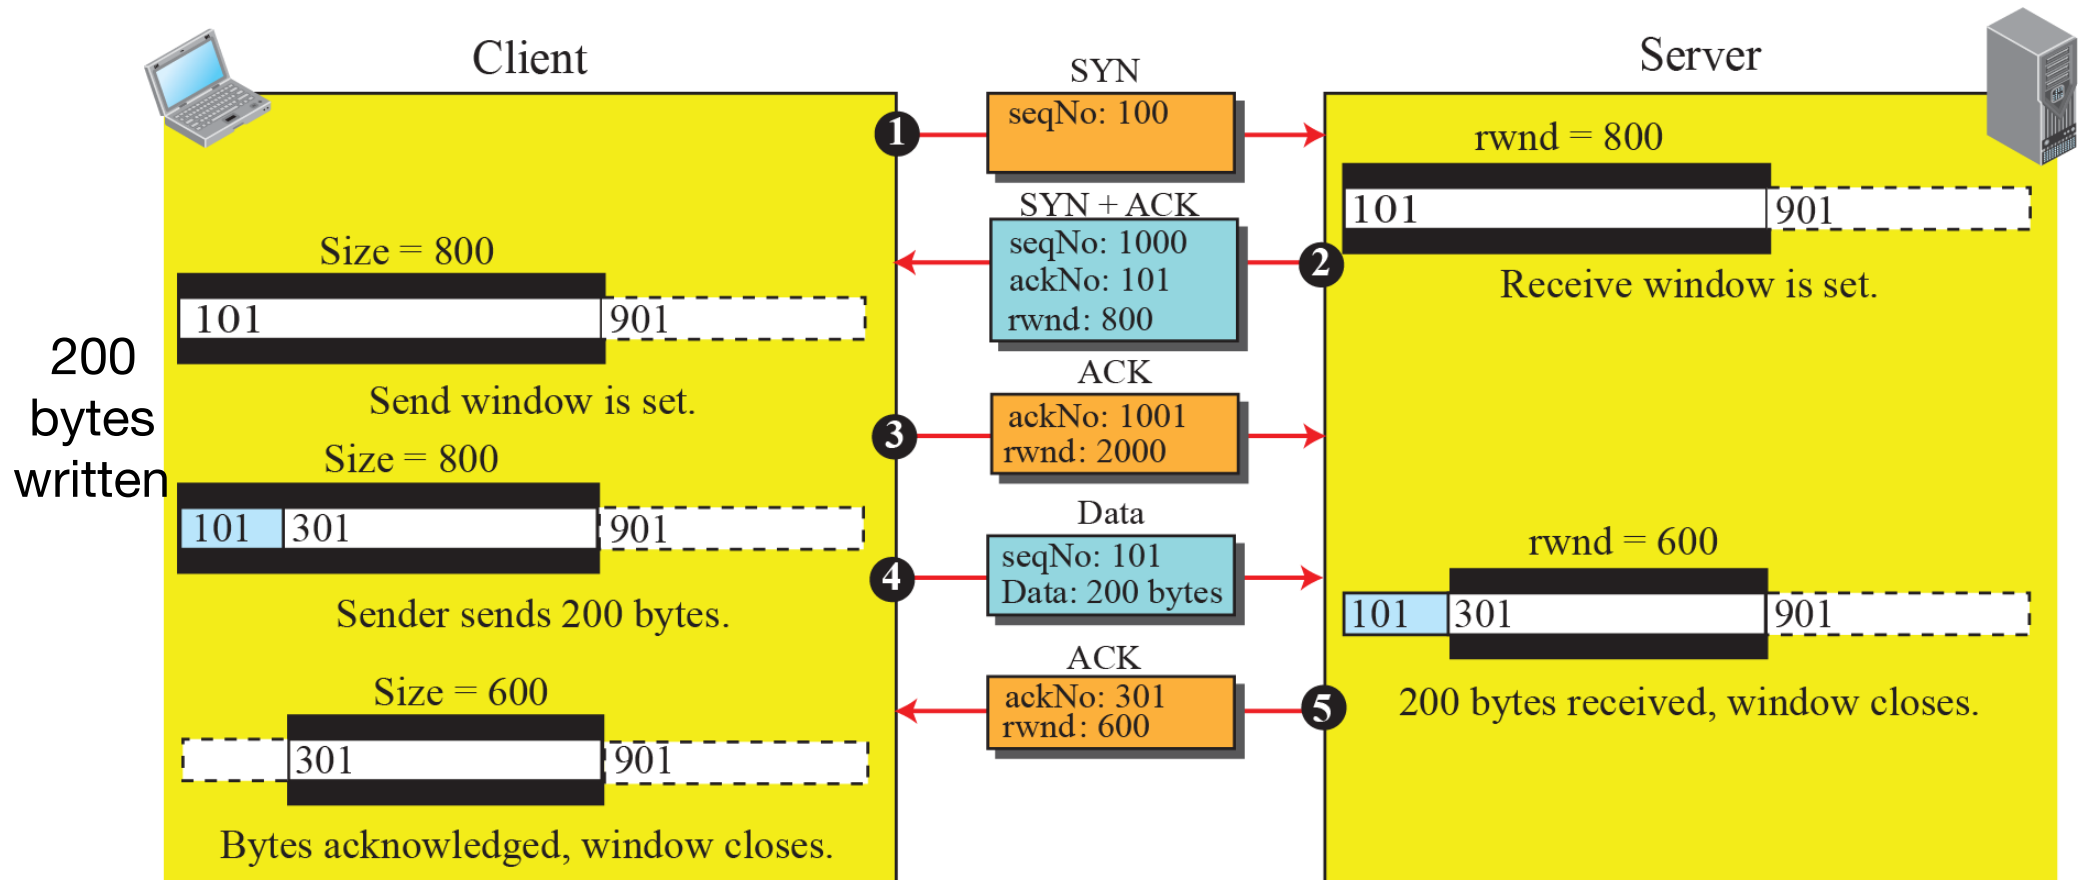
- ▶ Transport Layer Security

0020	f4 a7 01 bb f3 49 3b 86 93 a4 50 e4 64 1f 80 18	...I;. .P.d...
0030	04 33 54 4d 00 00 01 01 08 0a 3d 43 6a 33 1a 64	.3TM.... ..=Cj3.d
0040	d2 41 17 03 03 02 01 2e c9 93 fb 13 39 33 c8 62	.A..... .93.b
0050	52 2b 35 0b d5 12 b4 02 e0 a5 a7 ca b0 fa d7 0b	R+5.....
0060	5c 42 14 3a b2 b8 02 9b ee e3 bf a4 38 ff 54 42	\B.:.... .8.TB
0070	fb 31 cf f0 3d 82 9e 15 26 ee 5e 4f eb d3 23 de	.1..=... &^0..#.
0080	50 60 56 ca 2f b9 e5 0d 87 43 26 92 07 b4 5f 4c	P`V./... .C&..._L
0090	4c 9f 84 e2 2d 6b 19 80 0c 9d 78 e9 c2 31 7e a5	L...-k... ..x..1~.

TCP Connection Establishment: 3-Way Handshake



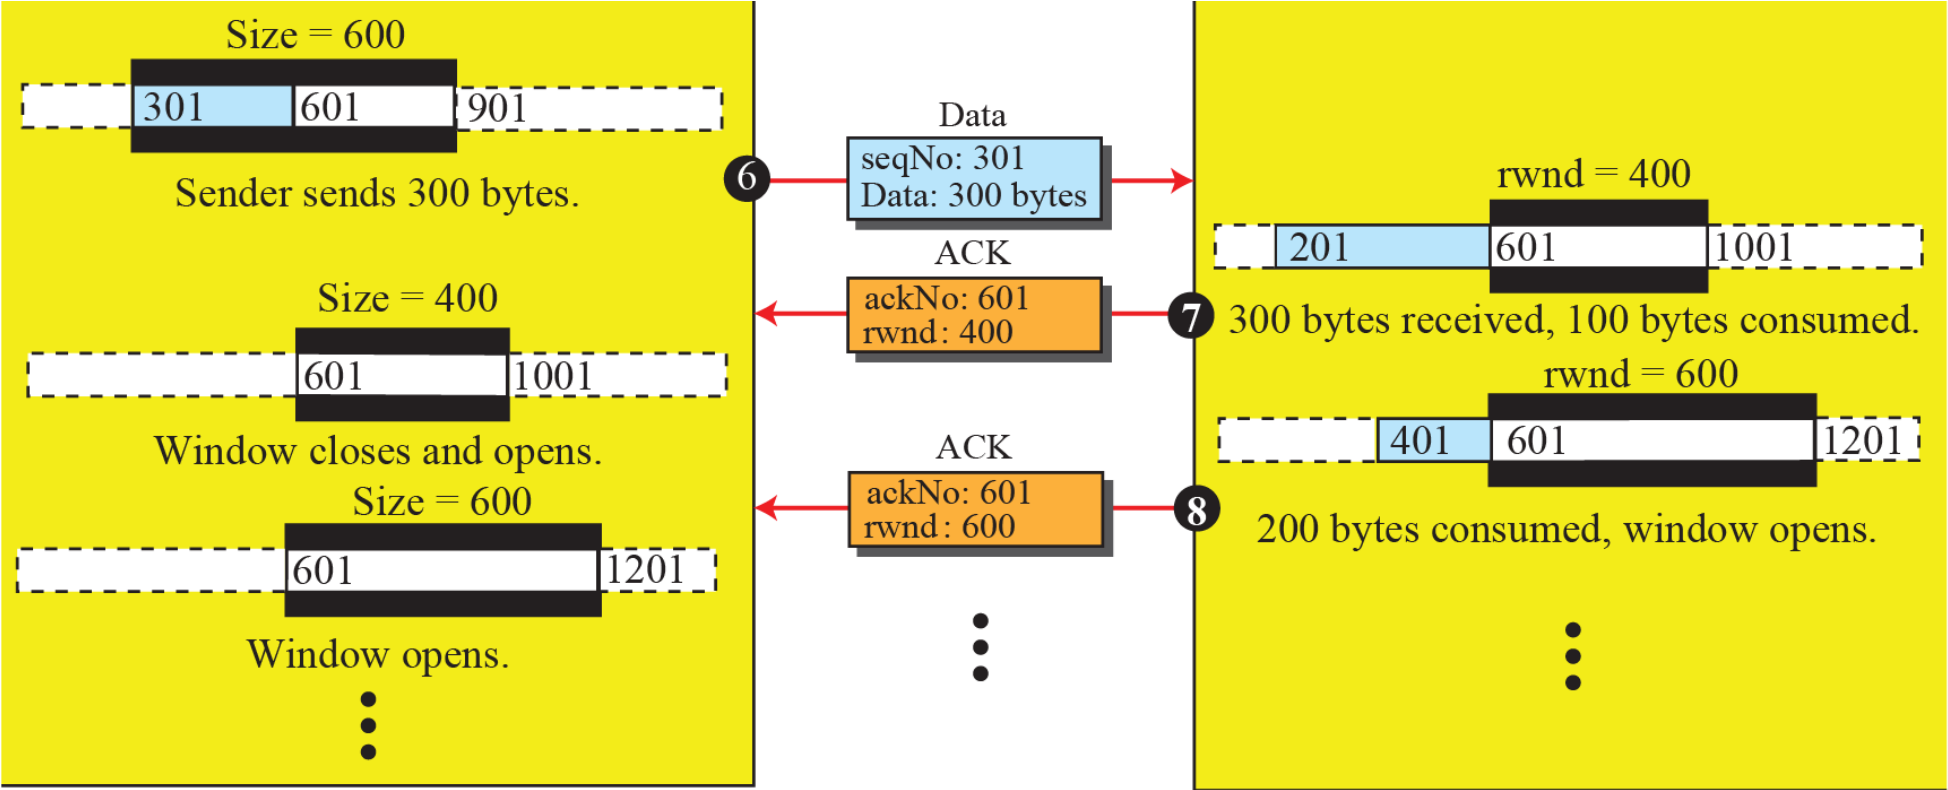
Rx buffer = 800 bytes



`rwnd` = Receive buffer space available

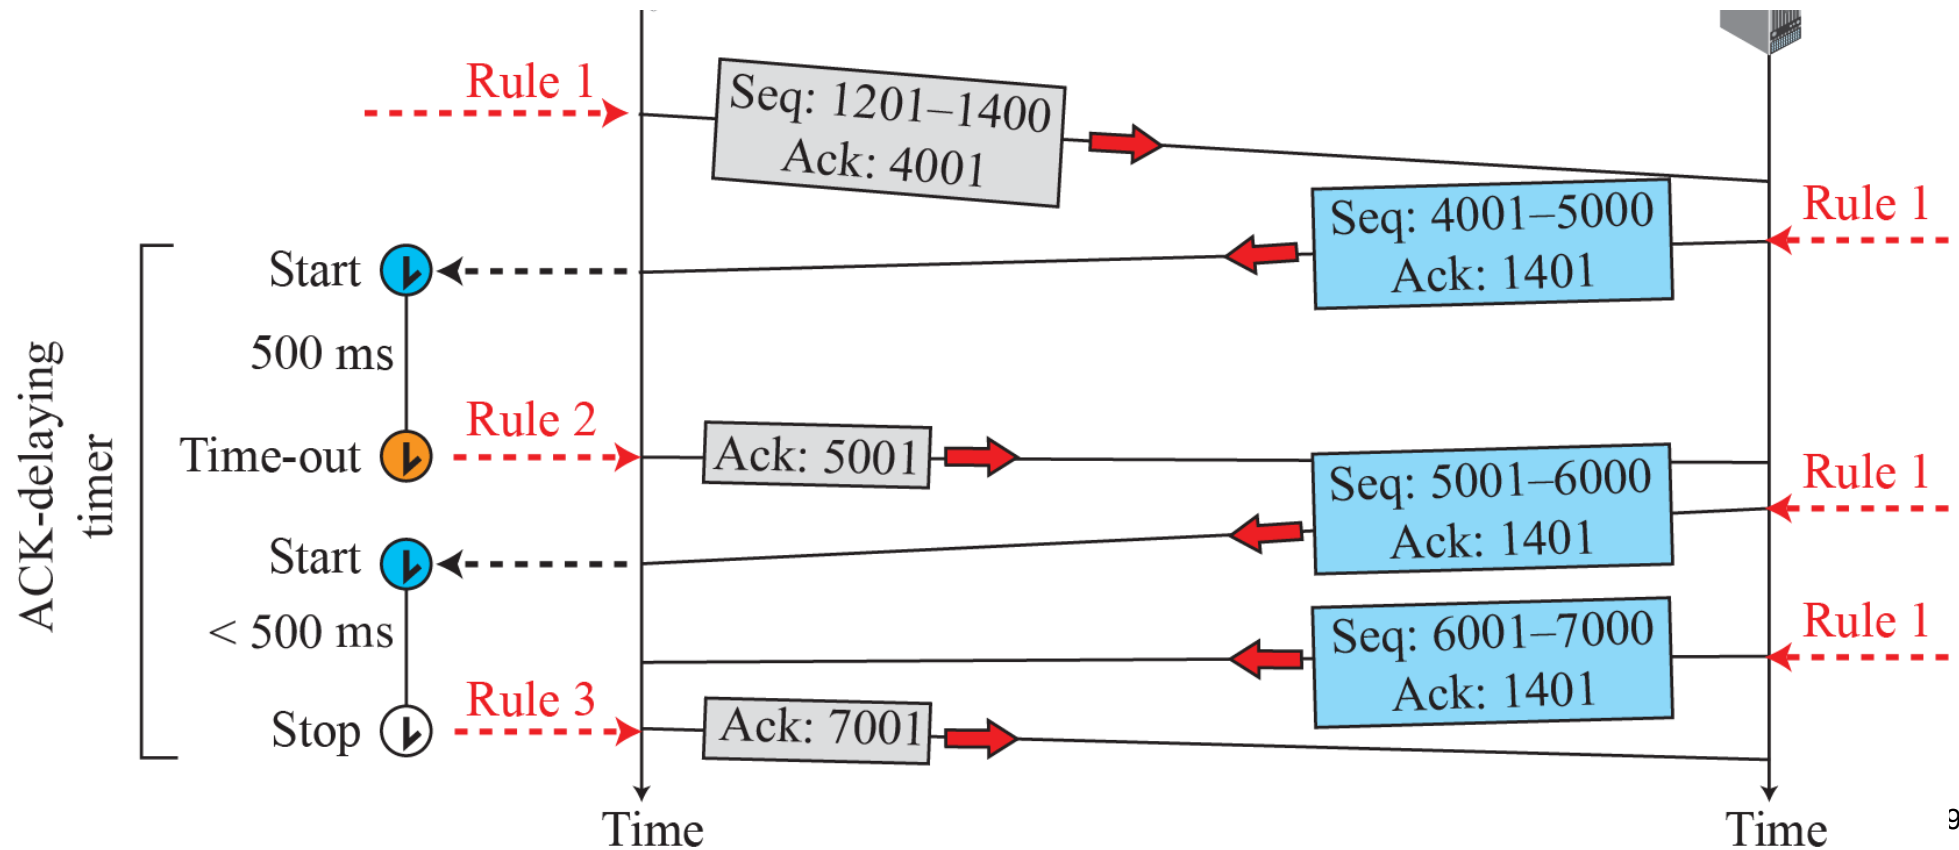
Receiving side keeps # bytes to send less than `rwnd`

Another 300 bytes written

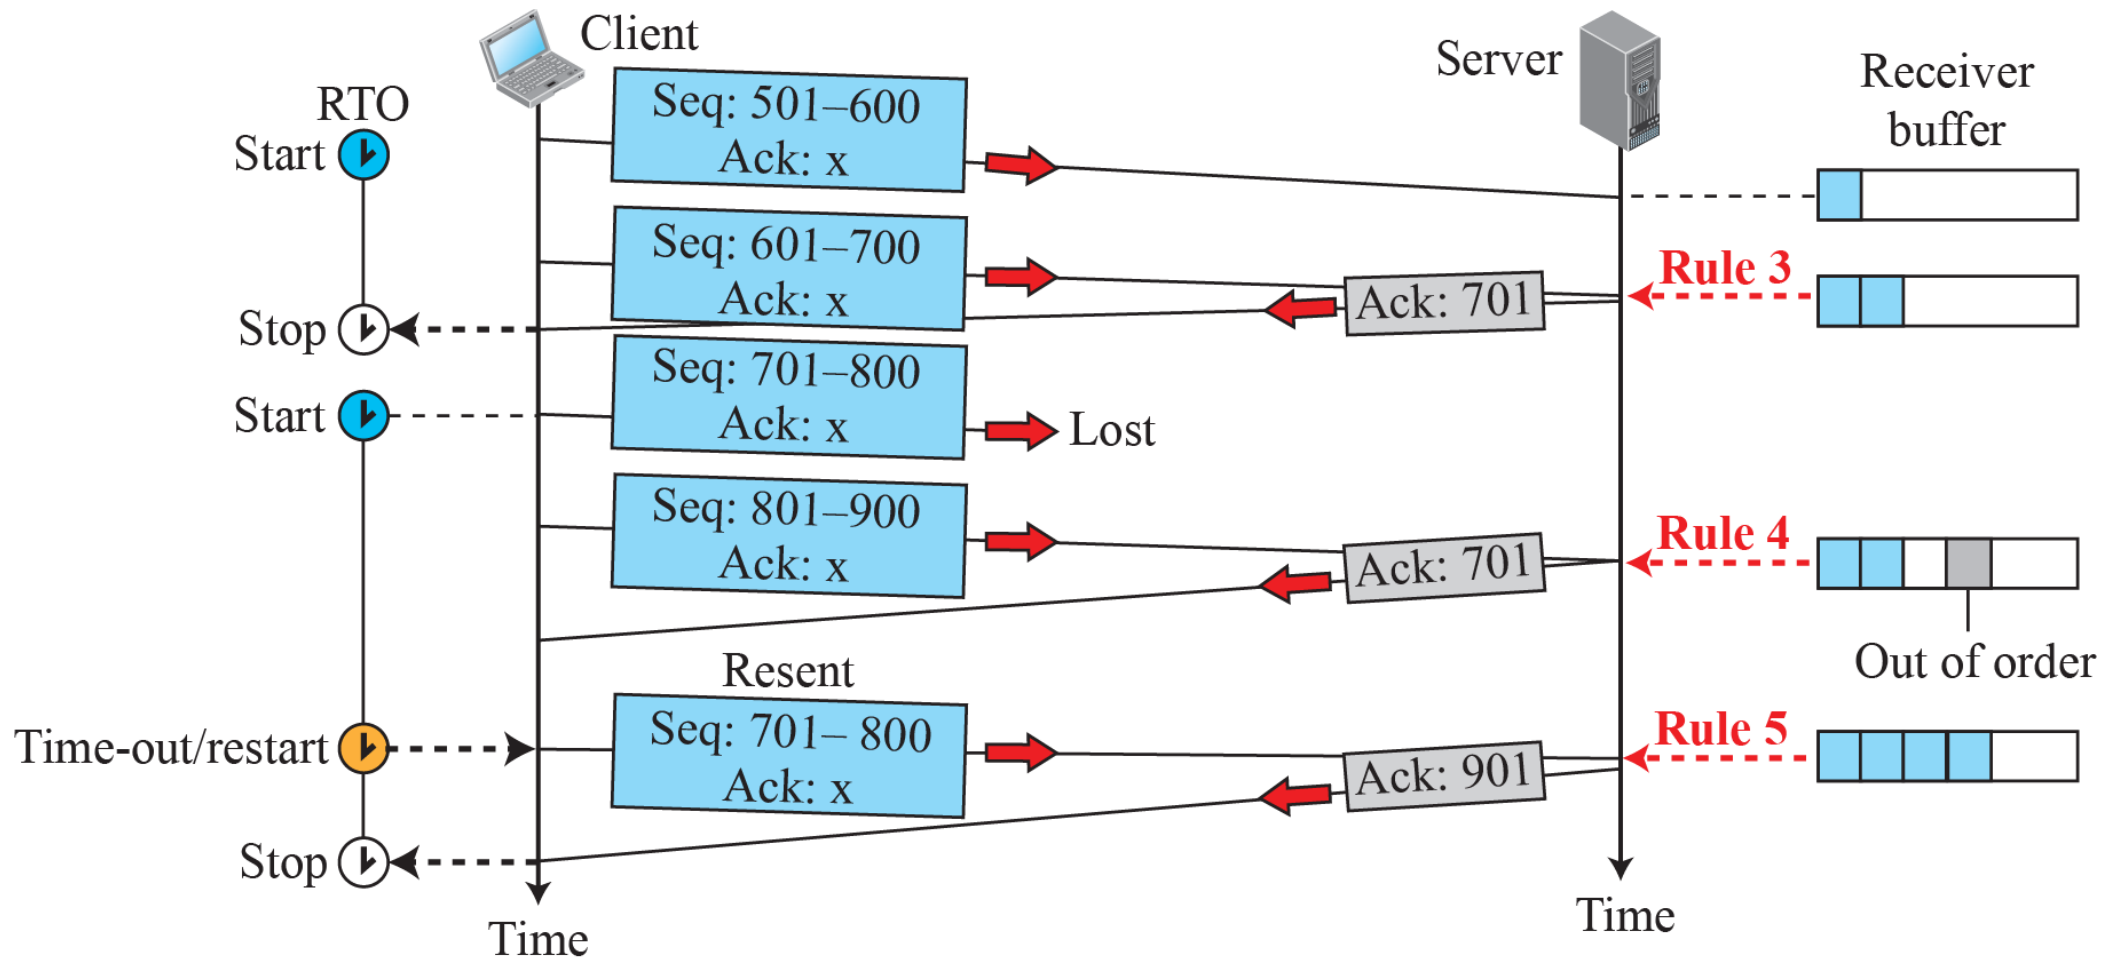


TCP Error Control -- Normal Operation

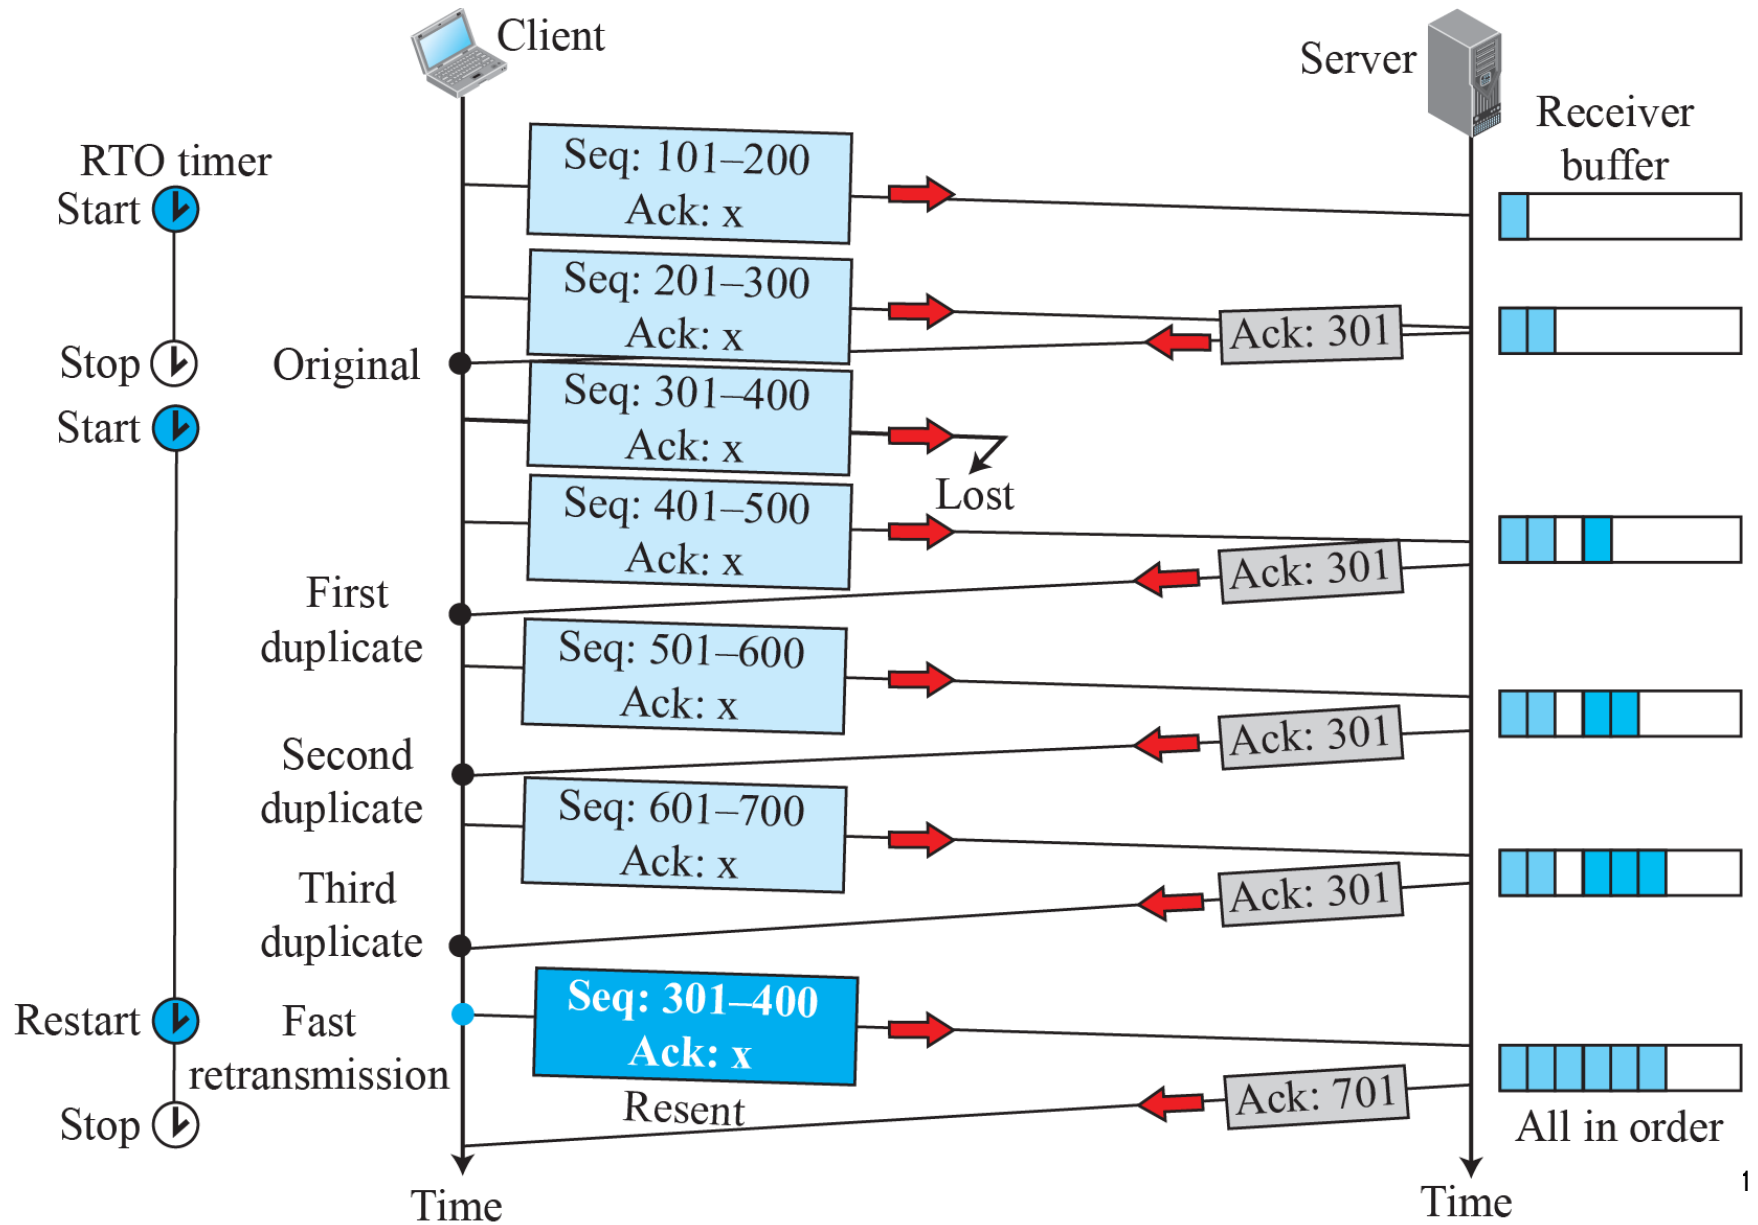
- See Acknowledgment rules in p.206 of Forouzan text
- Only one timer for the first outstanding segment.



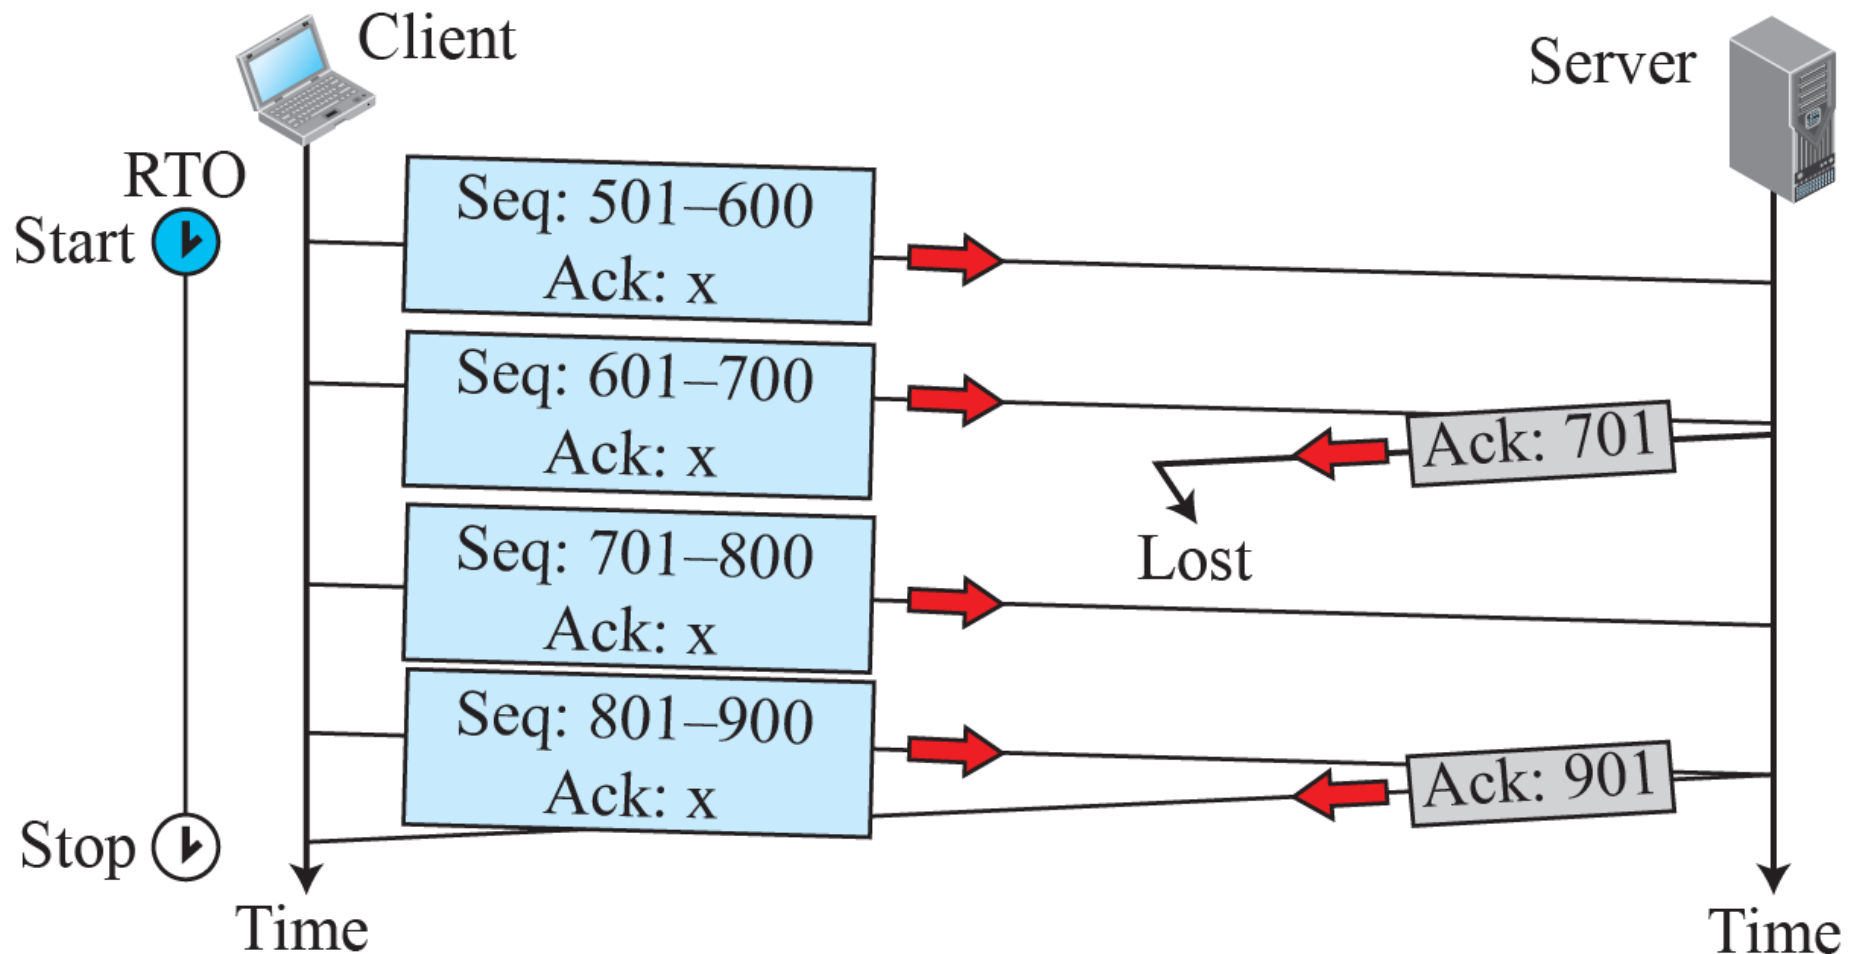
Lost Segment

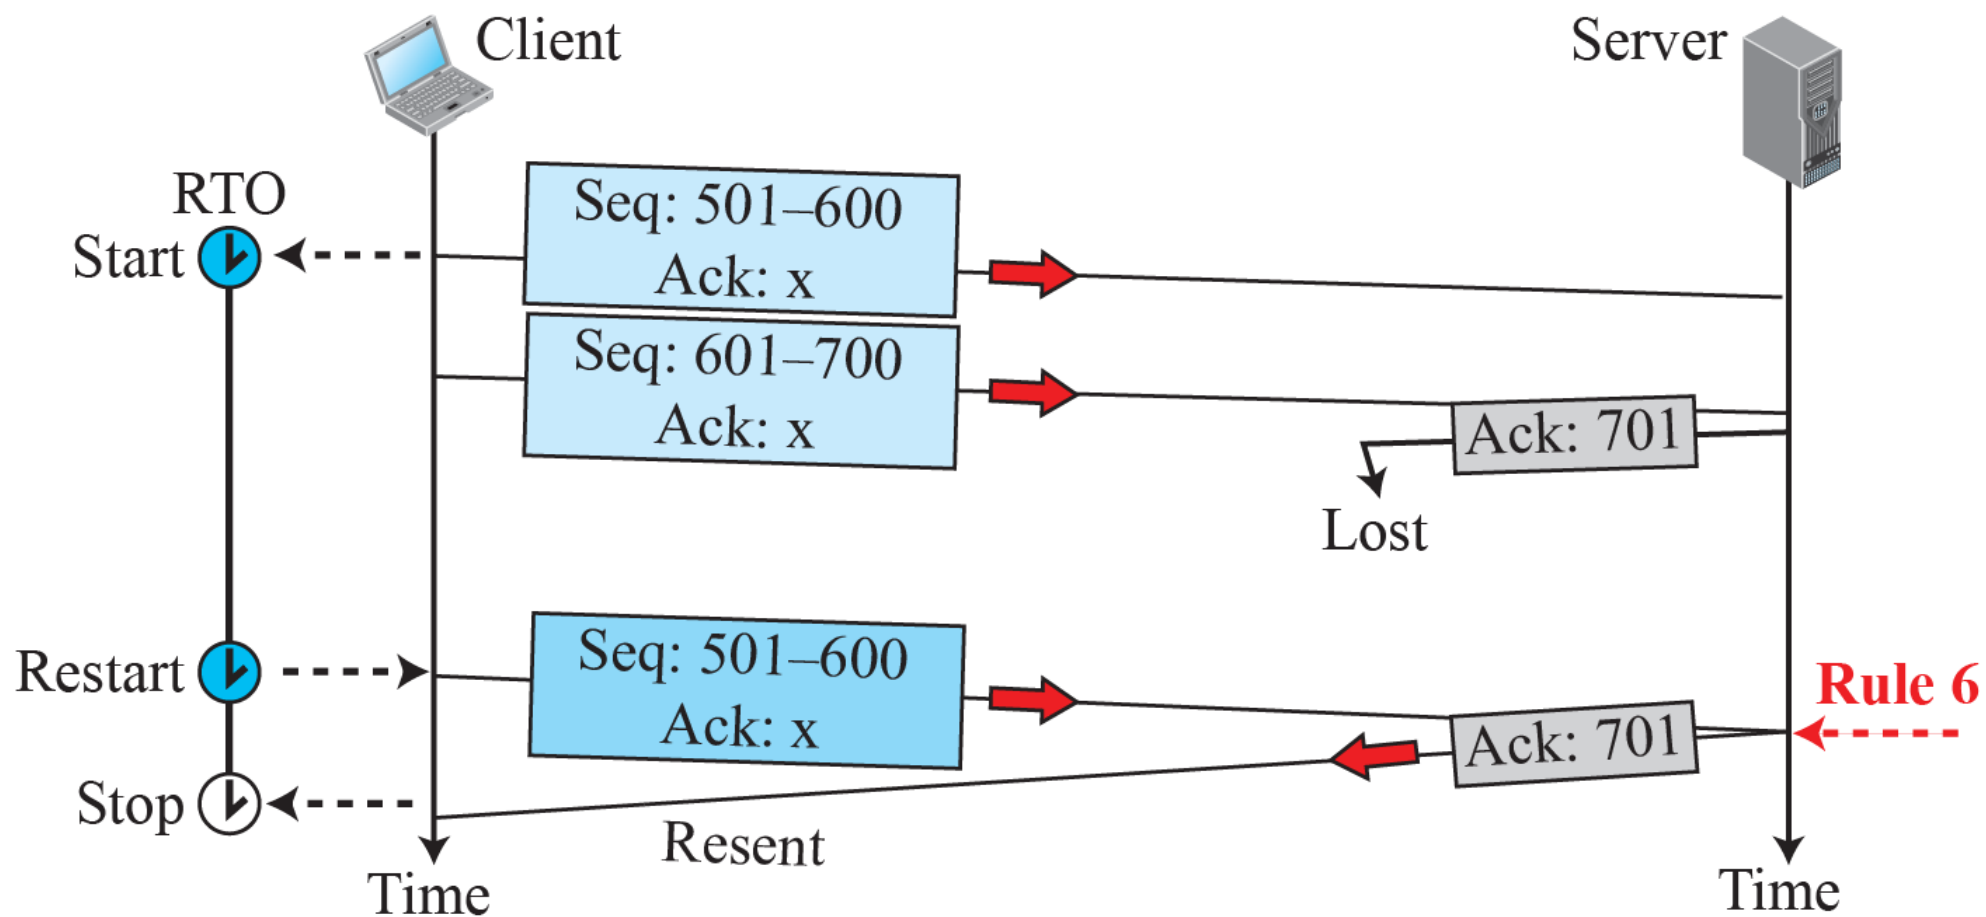


Fast Retransmission



Lost Acknowledgment





In Wireshark preference, go to protocols->tcp and checkbox the option 'Relative sequence number'

No.	Time	Source	Destination	Protocol	Length	Info
1549	4.816713	192.168.1.54	202.44.8.55	TCP	78	52942 → 80 [SYN] Seq=0 Win=65535 Len=0 MSS=1460 WS=32 TS...
1551	4.831000	202.44.8.55	192.168.1.54	TCP	78	80 → 52942 [SYN, ACK] Seq=0 Ack=1 Win=16384 Len=0 MSS=14...
1552	4.831071	192.168.1.54	202.44.8.55	TCP	66	52942 → 80 [ACK] Seq=1 Ack=1 Win=131744 Len=0 TSval=1022...
1553	4.831678	192.168.1.54	202.44.8.55	HTTP	955	GET / HTTP/1.1
1571	5.005340	192.168.1.54	202.44.8.55	TCP	955	[TCP Retransmission] 52942 → 80 [PSH, ACK] Seq=1 Ack=1 W...
1572	5.011084	202.44.8.55	192.168.1.54	TCP	66	80 → 52942 [ACK] Seq=1 Ack=890 Win=64646 Len=0 TSval=359...
1574	5.019968	202.44.8.55	192.168.1.54	TCP	66	[TCP Dup ACK 1572#1] 80 → 52942 [ACK] Seq=1 Ack=890 Win=...
1659	5.899773	202.44.8.55	192.168.1.54	HTTP	1514	HTTP/1.1 200 OK (text/html)
1660	5.899774	202.44.8.55	192.168.1.54	TCP	1514	80 → 52942 [ACK] Seq=1449 Ack=890 Win=64646 Len=1448 TSv...
1661	5.899858	192.168.1.54	202.44.8.55	TCP	66	52942 → 80 [ACK] Seq=890 Ack=2897 Win=128864 Len=0 TSval...
1665	5.924211	202.44.8.55	192.168.1.54	TCP	1514	80 → 52942 [ACK] Seq=2897 Ack=890 Win=64646 Len=1448 TSv...
1666	5.924215	202.44.8.55	192.168.1.54	TCP	1514	80 → 52942 [ACK] Seq=4345 Ack=890 Win=64646 Len=1448 TSv...
1667	5.924295	192.168.1.54	202.44.8.55	TCP	66	52942 → 80 [ACK] Seq=890 Ack=5793 Win=128160 Len=0 TSval...
1668	5.924357	192.168.1.54	202.44.8.55	TCP	66	[TCP Window Update] 52942 → 80 [ACK] Seq=890 Ack=5793 Wi...
1669	5.926060	202.44.8.55	192.168.1.54	TCP	1514	80 → 52942 [ACK] Seq=5793 Ack=890 Win=64646 Len=1448 TSv...

▼ Transmission Control Protocol, Src Port: 52942 (52942), Dst Port: 80 (80), Seq: 1, Ack: 1, Len: 889

Source Port: 52942

Destination Port: 80

[Stream index: 6]

[TCP Segment Len: 889]

Sequence number: 1 (relative sequence number)

[Next sequence number: 890 (relative sequence number)]

Acknowledgment number: 1 (relative ack number)

Header Length: 32 bytes


Self-Test

- How long your application needs to wait before it can send the first message to the destination after opening a TCP connection ?
 - A. One Round-Trip-Time (RTT)
 - B. One-and-a-half RTTs
 - C. Two RTTs
 - D. Four RTTs

Self-Test

- Suppose TCP continuously sends segments for your application. If one of the segments has lost but all the successive segments get through (before the RTO timer expires), approximately how long will it take for your application to receive the lost segment ?
 - A. One RTT
 - B. Two RTTs
 - C. Three RTTs
 - D. Four RTTs

TCP Connection Teardown



Exploiting TCP Connection Management

☐ TCP SYN scan

- Open a TCP connection to each port number one by one.
- SYN/ACK returned for open ports, RST for closed ports

☐ TCP FIN scan

- Send a FIN to each port number one by one.
- Nothing returned for open ports, RST for closed ports

☐ TCP SYN DoS attack

- Send SYNs to establish connections but doesn't return ACK.
- The victim runs out of memory and possibly crashes.

Activity I: Which Transport Protocol for My App ?

- Suppose you are to choose transport protocol (UDP and TCP) for your network application.
 - ◆ List and explain possible criteria that you will use to make the decision.
 - ◆ Can you prioritize the list ?

Conclusion

- End-to-end layer for process multiplexing, flow and error control, congestion control
 - Unreliable vs. Reliable transport service
 - Connectionless vs. Connection-oriented service
- TCP operations
 - Connection set up and teardown phases -- State initialization and resource management
 - Sliding window flow control and error control
 - High overhead for short transactions apps